

Mean rewards for different bandit algorithms, averaged 100 times
3 arms: Non-Stationary MAB, Bernoulli with $\Upsilon = 2$ break-points

