

Mean rewards for different bandit algorithms, averaged 100 times
 3 arms: Non-Stationary MAB, Bernoulli with $\Upsilon = 4$ break-points

